

First year project
Project 99: Who's Julia?

Group: 99a
Simon Lehmann Knudsen, simkn15
Sonni Hedelund Jensen, sonje15
Asbjørn Mansa Jensen, asjen15
FF501

May 30, 2016

1 Summary

Julia is a new programming language. Through tests and benchmarking it is found that * *

2 Preface

This report has been written in May 2016 as a first year project in the Bachelor Degree of Computer Science. We thank our supervisor Michal Kotrbčik for his help and time. In this project we have been doing benchmarks to find out if Julia is as good as the developers claim it to be.

Contents

1	Summary	I
2	Preface	I
3	Introduction	1
4	Julia	1
4.1	Syntax	2
4.2	Features	8
4.2.1	Dynamic typing	8
4.2.2	Multiple dispatch	9
4.2.3	User defined types	10
4.2.4	Garbage Collector	10
4.2.5	Built-in package manager	11
4.2.6	Lightweight green threading	11
4.2.7	Meta-programming and Macros	11
4.2.8	Implementing code from other languages	12
4.2.9	Extensible conversions and promotions for numeric and other types.	12
4.2.10	Designed for parallelism and distributed computation .	12
4.2.11	Automatic generation of specialized code for different argument types	12
4.2.12	Supports Unicode	12
4.2.13	Devectorized code is fast	13
5	Project Euler	13
6	Theory	13
6.1	Benchmarking	13
6.2	Profiling	14
6.3	Time measurements	14
7	Materials and methods	14
7.1	Benchmarking	14

8 Problems	16
8.1 Projecteuler 11	16
8.2 Projecteuler 116	19
8.3 Quicksort	20
9 Results	20
9.1 Benchmarking	20
10 Personal experience with Julia	20
10.1 Sonni	20
10.2 Simon	23
10.3 Asbjørn	23
11 Discussion	24
12 Conclusion	24
13 Perspective	24
14 Appendix	25
14.1 Process analysis	25
14.1.1 Projectmanagement, organizing workflow and organizing process of learning	25
14.1.2 Process of learning	25
14.1.3 Cooperation, difficulties in group and with supervisor	25
14.1.4 Conclusion	26
14.2 Copy of poster in A4	26

3 Introduction

Java, C, C++, C# and Python are on top ten of the most used programming languages. There are hundreds of languages, but not minding the hard competition new languages are still created with the thought of doing better. Julia is a new programming language, which has been developed to contain other languages best features.

The purpose of this report is to find out if the programming language Julia has a chance in the market and if it is worth changing to. The languages Java, C++, Python and Julia will be compared on syntax and time performance.

4 Julia

Julia is a programming language, which has been under development since 2009. First released in 2012 and the newest stable version of Julia is version 0.4.5. The language has been created because the developers wanted a language with all the features they like from different languages. The developers also wanted the language to be open source, which means everybody can read and modify the language. Julia shares many similarities with python. One of the idea was to make the language as simple, readable and easy to learn as possible. The language is made for high performance and scientific computations while still supporting general purpose programming. Programming in Julia can be done in the terminal like python with interactive mode:

(a) Julia

(b) Python

Writing longer programs in the terminal might not be the preferred method. The text editor Atom supports the Julia language. Atom has a package, **uber-juno**, which sets up Atom to act like an IDE, integrated development environment for Julia.

Make examples of syntax between all the languages here.

Figure 2: Declaring variables

```
1   for i = 1 : 10
2       println(i)
3   end
```

(a) Julia

```
1   for i in range(1, 11)
2       print i
```

(b) Python

```
1   for (int i = 1; i < 11; i++){
2       System.out.println(i);
3   }
```

(c) Java

```
1   for (int i = 1; i < 11; i++){
2       std::cout << i << "\n";
3   }
```

(d) C++

Figure 3: For-loops: Incrementing

```
1   for i = 10 : -1 : 1
2       println(i)
3   end
```

(a) Julia

```
1   for i in range(10, 0, -1)
2       print i
```

(b) Python

```
1   for (int i = 10; i > 0; i--){
2       System.out.println(i);
3   }
```

(c) Java

```
1   for (int i = 10; i > 0; i--){
2       std::cout << i << "\n";
3   }
```

(d) C++

Figure 4: For-loops: Decrementing


```
1   while i < 10
2       a -= 1
3   end
```

(a) Julia

```
1   while (i < 10):
2       a -= 1
```

(b) Python

```
1   while (i < 10){
2       i--;
3   }
```

(c) Java

```
1   while (a < 10){
2       i--;
3   }
```

(d) C++

Figure 5: while-loop

```
1   while i < 10
2       i -= 1
3   end
```

(a) Julia

```
1   while (i < 10):
2       i -= 1
```

(b) Python

```
1   while (i < 10){
2       i--;
3   }
```

(c) Java

```
1   while (a < 10){
2       i--;
3   }
```

(d) C++

Figure 6: while-loop

```
1     function hello(name)
2         println("Hello $name")
3     end
```

(a) Julia

```
1     def hello(name):
2         print "Hello", name
```

(b) Python

```
1     public static void hello(String name){
2         System.out.println("Hello " + name);
3     }
```

(c) Java

```
1     void hello(std::string name)
2     {
3         std::cout << "Hello " << name << std
4             ::endl;
```

(d) C++

Figure 7: Declaring functions

```
1 array = fill(0, 100)
```

(a) Julia

```
1 array = [0 for col in range(100)]
```

(b) Python

```
1 int[] array = new int[100];
2 Arrays.fill(array, 0);
```

(c) Java

```
1 std::array<int,100> array;
2 array.fill(0);
```

(d) C++

Figure 8: Filling array

4.2 Features

4.2.1 Dynamic typing

Dynamic typing means that type checks are mostly performed at runtime - opposed to static typing, which makes the type checks at compile time. Dynamic typing allows the programmer to skip the type declaration and let the dynamic type system handle it. In Julia it is possible to specify types, but mostly is not necessary. Julia also makes decisions about how much memory to allocate for variables and the assigned memory size will not change later, not even if a big value, that exceeds the memory of the variable, is assigned to it. The memory size can be decreased and this will be addressed later under the section Garbage Collector. However, there is a range of memory sizes to chose from when initializing a variable. Julia will by default assign 32 bits to an integer or a float if no more memory is needed at initialization, but it is possible to manually chose one of the following:

- Signed / unsigned integers of: 8, 16, 32, 64 and 128 bits.
- Floating points of: 16, 32 and 64 bits.
- Boolean – 8 bits
- Char – 32 bits

To initialize a variable with a given memory size:

```
1 x = Int8(10)
```

But if the variable is later assigned to a new value. Julia will automatically allocate 32 bits for the variable and more if needed:

```
1 x = Int8(10)
2 typeof(x) #Int8
3 x = 4
4 typeof(x) #Int32
```

If no certain memory size is needed but only a specific type then abstract types will come in handy. An abstract type is just some form of generalization of a concrete type. For example 16, 32 and 64 bit floating points are of type `AbstractFloat`. This is especially useful when the programmer do not know how much memory a certain variable may need. Too little memory will result in a bug and too much memory is a waste.

4.2.2 Multiple dispatch

Multiple dispatch is used to determine, which function to call by the type of one or more of the parsed argument(s). This is useful for example when two functions with the same name are declared:

```
1 function a(arg1::Int8)
2     println("Int")
3 end
4
5 function a(arg1::Float16)
6     println("Float")
7 end
```

If a variable is declared with a type of `Int8` and parsed to the function `a(...)`, then it will print "Int", and if the variable is declared as `Float16` it will print "Float". It is also possible to use abstract types here, so `arg1::Float16` and `arg1::Int8` could be changed to `arg1::AbstractFloat` and `arg1::AbstractInt` – this will make the function 'a' accept any kinds of floating points and integers.

4.2.3 User defined types

User defined types, also known as composite types gives the option to define new types in Julia. A composite type in Julia may look something like the following:

```
1 type person
2     age::Int16
3     name #a variable, which gets the assigned default type ::Any
4 end
```

If a variable in a composite type is not specified with any type, the default is `::Any`, which accepts any kinds of types. To initialize a composite type, think of it like an object. `Person = person(21, "Carl")`, and the values can be changed with `"Person.age = 25"`. In the example above, the type `person` can now be used as any other of the built in types in Julia.

```
1 function a(arg1::person)
2     println(arg1.age)
3 end
4
5 Person = person(21, "Carl")
6
7 a(Person) #Will print 21
```

The developers of Julia state that user defined types are as fast and compact as built in types.

4.2.4 Garbage Collector

Julia uses a garbage collector to automatically free the memory when needed. There is no guarantee when the garbage collector will run, but it is possible to force garbage collection with a function call: `gc()`. One thing to keep in mind is that once a name is defined in Julia, it will always be present till termination. The garbage collector does not free the memory of unreachable object but rather reallocates memory for objects when memory size has changed. So the only way the garbage collector can free memory is when an object has been reduced in size, for example:

```
1 A = rand(float32, 10000, 10000)
2 A = 0
```

```
3 gc()
```

This code will generate a 10.000x10.000 matrix filled with random 32 bit floating points, which consumes a bit of memory. **A** will be set to 0 and will keep the memory size of the 10.000x10.000 matrix until garbage collection is done either manually as in this example, or somewhere in the future when it is done automatically.

4.2.5 Built-in package manager

Julia comes with a built in package manager, which keeps track of packages that needs to be included in the program to run. It is not necessary to state, which packages needs to be included. For example, when the user calls the sort function, **sort()**, no package include statement or **math.sort()** is needed – this is done automatically by the package manager.

4.2.6 Lightweight green threading

A thread is lightweight when it shares address space with other threads opposed to heavyweight threads, which has its own address space. If threads share the same address space the communication between them is faster and much simpler. The communication between heavyweight threads have to go through pipes or sockets.

Green threads are threads that are scheduled not by the operating system but rather by the runtime library or virtual machine. Green threads does not have to relay on the operating system and can be controlled much better – which means the threads are in user space and not kernel space, these terms will be adressed later in **Benchmarking**. As of Julia 0.4.5 multithreading is not available but in the unstable version 0.5.0 an experimental support of multithreading is available, therefore threading will be addressed no further.

4.2.7 Meta-programming and Macros

Meta-programming is a way to write programs in programs and let them use program code as data. In Julia meta-programming can be done by defining macros. For example, the **@time** macro was used extensively in the project. The **@time** is in the standard Julia library. The **@time** takes a

function as an argument and sets a timer in the top of the code from the argument and stops the timer in the bottom and prints the time passed and memory allocated. Meta-programming and macros are known from Lisp, a programming language, and have been used in early AI research.

4.2.8 Implementing code from other languages

Since Julia is a newer language there are not many written libraries yet. Julia has an import feature for both Python and C libraries to avoid this disadvantage. Import the library PyCall with the "using" operation to use python and import Python libraries via a macro `@pyimport`. To use C libraries or coding, simply run the function `ccall()`. The programming language C is a well known and used language. Many systems are based on C, which makes C an advantage for hardware programming. Another feature is to execute shell commands. Execute shell commands using `run()`. An example would be running `run('echo Hello World')`, which would return the output "Hello World".

4.2.9 Extensible conversions and promotions for numeric and other types.

4.2.10 Designed for parallelism and distributed computation

Julia claims to be designed for parallelism and distributed computation. Parallel computing is dividing a problem into smaller problems and solve the smaller problems at the same time. Distributed computing is solving a problem using a network of computers working together to get to the solution. Julia have been trying to make this easier to achieve with different macros.

4.2.11 Automatic generation of specialized code for different argument types

4.2.12 Supports Unicode

The Unicode support globalize the language. Unicode is an encoding system for characters. Unicode supports a large range of characters and is designed to make it easier to read and write non-latin characters.

4.2.13 Devectorized code is fast

Vectorization of code is, for example, when a for loop is rewritten to sequential operations. Figure 9 shows an example of a non-vectorized for loop and the vectorized version. Vectorized code is usually faster. Julia claims that there is no reason to vectorize as the performance stays the same.

```
1      a = 1, 2, 3, 4
2      b = 2, 3, 4, 5
3      c is empty
4      for i = 1 to 4
5      begin
6          c[i] = a[i] + b[i]
7      end
```

(a) Devectorized code

```
1      a = [1, 2, 3, 4]
2      b = [2, 3, 4, 5]
3      c is empty
4      c[1] = a[1] + b[1]
5      c[2] = a[2] + b[2]
6      c[3] = a[3] + b[3]
7      c[4] = a[4] + b[4]
```

(b) Vectorized code

Figure 9: Vectorization

5 Project Euler

Project Euler is a website with a huge database of mathematical and programming challenges. In this report Project Euler has been used to find challenges for testing the different programming languages.

6 Theory

6.1 Benchmarking

Benchmarking is measuring the quality of a product. In this report we will be measuring the time it takes for a programming language to execute an

algorithm.

6.2 Profiling

Profiling is a way to analyse code. For example, the analysis could be measuring the memory usage, time and function calls. The analysis shows, which part of the code took the longest time or used up all memory. Profiling is often used for optimization.

6.3 Time measurements

There are three types of time measurements. Wall, user and system time. Wall time is the time it takes in real time and wall time is affected if other program are running. System time is the cpu time used inside kernel. The time used inside kernel is for memory allocation, file reading, file writing and other thing going on outside the cpu and memory. User time is the time used calculating and comparing etc. In this report the sum of system and user time is used, because the system and user time are not affected by heavy loads on the machine and only shows the actual time taken for the Computer to solve the problem.

7 Materials and methods

7.1 Benchmarking

For the benchmarking we used: Lenovo W530: Intel Core i5-3320M 2.60GHz x 4, 8GB RAM, Ubuntu 16.04 64-bit, Windows 10 64-bit, Windows XP 32-bit MacBook: Intel Core i5 1,4GHz, 8GB RAM, OSX El Capitan 10.11.5 MacBook: Julia 0.4.5 Python 2.7.11+ C++11 Java 1.8.0_92 All time measurements are measured using the shell 'time' command that returns the three types of time taken for the algorithm to run. The problems are scaled in three different sizes and run in each language several times. A tool created for this project gather the data and calculates the average and median running time.

The benchmarking will be about the running time of each individual problem. When looking at running time the most common terms are wall time (real time) and CPU time. Wall time is the time it took the program to terminate, from start to finish. CPU time is how much time the program used

on the CPU. The time difference between wall time and CPU time comes from the fact that the program being benchmarked might get interrupted by the operating system, because other tasks need to be handled. These tasks could be other programs, or something that the OS needs to get done. The processing time of those tasks will be included in the wall time but excluded in the CPU time. To get the running times in Unix based system, the time command is available from the terminal:

```
1 time <COMMAND>
```

Where <COMMAND> can be any terminal command. The time command will return three different times.

```
real    0m0.627s
user    0m0.462s
sys     0m0.147s
```

Figure 10: Output of time

- Real: This is the wall time mentioned above.
- User: User time is how much CPU time is spent in user mode.
- Sys: System time is how much CPU time is spent in kernel mode.

The difference between user- and kernel mode is the following:

In most memory protected system there are some locked operations for security reasons. E.g. allocation of memory or accessing hardware requires kernel mode and cannot be done in user mode. Processes like malloc and reading / writing to files will somewhere in the process send a request to the kernel to do certain things and this is counted as system time. This does not necessarily mean that all time is spent in kernel mode when reading a file, since the only time that is spent in kernel mode is accessing the requested data in memory the rest of the operation is counted towards user time.

So to get much times is spent on the CPU, User time + system is a good estimation. To get an even better indication of the real CPU time, multiple test are done and the average is calculated. The average running time is considered since there can be fluctuations in run times for every test.

Another way to measure time spent on the CPU is to use the built in libraries for example `System.currentTimeMillis()` in Java. This is especially useful for profiling and figuring out, which part of a program is time consuming. Benchmarking will be done by comparing CPU time of the same problem in different languages. The problems will be scaled to have smaller and larger inputs. However, in some situations a large input can cause memory problems, which will be addressed in the specific problem. The memory problems is unfortunately making an upper bound for how large data is possible to pass for certain problems.

The algorithm used for a given problem is the same used in all compared languages. Extended tests will be included for, which the algorithm is optimized for a given language.

8 Problems

8.1 Projecteuler 11

In the 20x20 grid below, four numbers along a diagonal line have been marked in red. The product of these numbers is $26 \times 63 \times 78 \times 14 = 1788696$. What is the greatest product of four adjacent numbers in the same direction (up, down, left, right, or diagonally) in the 20x20 grid?

```

4851 3238 1738 4414 2453 3781
1058 2209 2485 2302 1601 4652
4791 3553 706 2520 4236 3842
1210 4707 2305 497 4336 3310
2611 3945 2743 1781 4386 3920
2476 1321 1870 471 1313 3326

```

The program runs from command-line/terminal with command line arguments in order to be able to test on different inputs and amount of adjacent numbers multiplied. To run the Julia version from the terminal:

Julia euler11.jl 100 4.

This would make the program run with a matrix of size 100 and calculating product of 4 adjacent numbers. A matrix generator was created, which can be found in appendix, to produce the matrix data. The matrix generator

requires one argument, which is the size of the matrix. So if 500 is passed to the matrix generator, a file named 500mat.txt will be created containing 500x500 digits ranging from 100-9999.

The algorithm first reads through the input file and creates a matrix. To calculate products in all directions needed in the problem, the algorithm goes through the matrix a total of 3 times. A nested for loop is needed to go through a matrix, lines 1-2 at figure 11. The outer loop iterates through the rows, and the inner loop iterates through the columns. Figure 11 calculates the horizontal and vertical directions, figure 12 diagonally from left to right(downwards) and figure 13 going diagonally from right to left(downwards). The algorithm starts in the most upper left cell, and iterates through all cells in the matrix. Lets have the first cell as (1,1), first number is row(**i**) and second number is column(**j**). Looking at figure 11, **matLength** is the size of the matrix, size = 100 would mean a matrix of size $100 \cdot 100$. **numProd** is the number of adjacent numbers multiplied together. For matrix with size 100, and multiplying four adjacent numbers, the algorithm would do the following in figure 11:

Starting at cell (1,1) to the end, which is cell (100, 100 - 4).

Line 5-7: Loops over the adjacent cells and multiplies the numbers.

Line 8-10: Sets current product to max product, if current is larger than previously max product.

```

1  for i = 1 : matLength
2      for j = 1 : matLength - numProd
3          #right/left
4          prod = Int64(1)
5          for k = 0 : numProd - 1
6              prod *= mat[i, j + k]
7          end
8          if prod > maxProd
9              maxProd = prod
10         end
11         #up/down
12         prod = Int64(1)
13         for k = 0 : numProd - 1
14             prod *= mat[j + k, i]
15         end
16         if prod > maxProd
17             maxProd = prod

```

Figure 11: Horizontal and vertical

Figure 12 shows the loop that iterates over the matrix and calculating the product of the diagonal going downwards from left to right. Figure 13 shows the diagonal product going upwards from left to right. The loop starts in the first row, and starting column is the last, not the first like the previously loops.

```

1  #diagonal left->right
2  for i = 1 : matLength - numProd
3      for j = 1 : matLength - numProd
4          prod = Int64(1)
5          for k = 0 : numProd - 1
6              prod *= mat[i + k, j + k]
7          end
8          if prod > maxProd
9              maxProd = prod

```

Figure 12: Diagonal downwards left to right

```


1  #diagonal right->left
2  for i = 1 : matLength - numProd
3      for j = matLength : -1 : numProd
4          prod = Int64(1)
5          for k = 0 : numProd - 1
6              prod *= mat[i + k, j - k]
7          end
8          if prod > maxProd
9              maxProd = prod

```

Figure 13: Diagonal upwards left to right

8.2 Projecteuler 116

Description: A row of five black square tiles is to have a number of its tiles replaced with coloured oblong tiles from red(length two), green(length three), or blue(length four). If red tiles are chosen there are exactly seven ways. If green tiles are chosen there are three ways. And if blue tiles are chosen there are two ways. Figure 14 is a visualization of how the tiles can be lain. Assuming that colours cannot be mixed there are $7 + 3 + 2 = 12$ ways of replacing the black tiles in a row measuring five units in length. How many different ways can the black tiles in a row measuring fifty units in length be replaced if colours cannot be mixed and at least one coloured tile must be used?



(a) Red tiles



(b) Green tiles



(c) Blue tiles

Figure 14: Projecteuler: 116

8.3 Quicksort

Quicksort applies the divide-and-conquer paradigm to sort numbers. Divide-and-conquer has three steps. **Divide** the problem into a number of subproblems that are smaller instances of the same problem. **Conquer** the subproblems by solving them recursively. If the subproblems sizes are small enough, however, just solve the subproblems in a straightforward manner. **Combine** the solutions to the subproblems into the solution for the original problem [1, chapter 4, page 65]. Description of quicksort: [1, chapter 7, page 170-171]

Divide: Partition (rearrange) the array $A[p..r]$ into two (possibly empty) subarrays $A[p..q-1]$ and $A[q+1..r]$ such that each element of $A[p..q-1]$ is less than or equal to $A[q]$, which is, in turn, less than or equal to each element of $A[q+1..r]$. Compute the index q as part of this partitioning procedure.

Conquer: Sort the two subarrays $A[p..q-1]$ and $A[q+1..r]$ by recursive calls to quicksort.

Combine: Because the subarrays are already sorted, no work is needed to combine them: the entire array $A[p..r]$ is now sorted.

Figure 15 shows a visualization of quicksort.

9 Results

9.1 Benchmarking

10 Personal experience with Julia

10.1 Sonni

Some of the design choices, in Julia, can be really frustrating for someone who is adapting to Julia. The indexes starts at 1 - this could be debatable because of the statement that Julia is used in a lot of scientific computations and many programming languages made for scientific computations starts at index 1. But as a programmer who is used to indexes starting at 0, this can cause a lot of unnecessary bugs. The developers of Julia also made a weird design choice with the for-loop syntax. A incrementing for loop looks like:

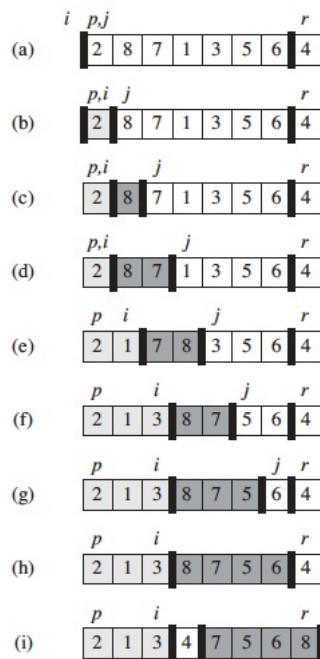


Figure 15: Quicksort

```

1 Julia:
2 for i = 1 : n
3     Some code ...
4 end
5
6 General languages:
7 for i = 1; i < n; i++
8     Some code ...

```

This is straightforward like many other languages, but take a look at the decrementing for-loop:

```

1 Julia:
2 for i = n : -1 : 0
3     Some code ...
4 end
5
6 General languages:
7 for i = n; i > 0; i--

```

8 | Some code

In Julia the exit condition and incrementation / decrementation are swapped and there is no consistency. The decrementing for-loops both start at `n`, and stops at 0. This kind of syntax can again cause bugs, and is not that easy to get eyes on.

The documentation of Julia is far from being great, which is a huge bump on the road to learn Julia. It feels far from complete. The documentation do not explain much but instead gives some examples. Essential information are often lacking, e.g. what types is needed to parse for a given function. It might point out that these arguments are needed, but lacking to describe those arguments. An example from the documentation is the function `rand()`, which is used to generate random numbers:

```
1 rand( [ rng ] [ ,S ] [ , dims... ] )
2
3     Pick a random element or array of random elements from the set
      of values specified by S; S can be:
4     - an indexable collection (for example 1:n or ['x','y','z']), or
5     - a type: the set of values to pick from is then equivalent to
      typemin(S):typemax(S) for integers (this is not applicable to
      BigInt), and to [0,1) for floating point numbers;
6     S defaults to Float64.
```

The argument `S` is explained, but does not mention anything about `rng` or `dims`.

Another example is how to access data within an array at a specific index. This should be one of the first things in the Array section, but is first described after about two pages of text. A person with experience in programming might know that to access an element in an array with an index is this simple line of code:

```
1 a[index]
```

But this isn't so obvious for new programmers. It is clear that something has to be done with the documentation.

Another downside to Julia is its consistency. The developers has tried to

remove some of the burden from the programmer with dynamic typing, package manager system etc., but with memory allocation and garbage collection it does not make much sense. The garbage collector does not free memory from unreachable object. Sizes of variables cannot be changed after declaration, and is not automatically increased if needed. If the programmer assign for example 8 bit memory to a variable and later change the value of that variable, then julia automatically assigns 32 bit or more if needed to that variable even if the 8 bits were enough.

10.2 Simon

I personally see myself the one in the group with the least programming experience, and might need more time learning Julia. I was not learning as fast as the others. I only know a few languages, and as many know, the more languages you know, the easier it gets learning a new language. The documentation of Julia was not satisfying enough to me. I learn faster if I can visually see how a function is used, rather than reading a wall of text. The visualizations on the current documentation was not fulfilling my needs. The documentation also misses to explain crucial information to a function, like what a given argument is presenting. For that reason I often found myself looking through the documentation of Python, to read about a function, since the syntax is so similar. Nevertheless, for programming I keep finding Julia more satisfying as I learn more, and for future programming I will definitely use Julia over Python.

10.3 Asbjørn

I found the syntax a bit confusing at first. I am not used to Python and I missed code structure and forced typing, in Julia and Python it looks like the code is hovering. After some time I got used to it and I found it very easy to learn. It makes it a little hard to find answers with the documentation and community being a bit thinner than with other languages. I think the 1-indexing is a bit complicated because I had never thought any language would use this, but I found that a lot of languages actually use this, they just are not that popular.

11 Discussion

Pros and Cons with good arguments

12 Conclusion

Is Julia a worthy competitor? Is it worth switching to Julia? Would we recommend Julia to anyone?

It can be difficult to make any conclusions by the benchmarking but it may give an idea of how Julia performs, compared to Java, Python and C++. One thing to keep in mind is that as of writing this report, Julia is in version 0.4.5 - it has yet to reach version 1.0.0 while the other languages are much older than Julia. This could mean that Julia might get even faster in the future.

13 Perspective

Articles to perspective:

<http://www.evanmiller.org/why-im-betting-on-julia.html>

High-Performance JIT Compiler chart at <http://julialang.org/>

<http://www.nowozin.net/sebastian/blog/the-julia-language-for-scientific-computing.html>

<http://radar.oreilly.com/2013/10/julias-role-in-data-science.html>

* Used too much times searching for articles.. *

14 Appendix

14.1 Process analysis

14.1.1 Projectmanagement, organizing workflow and organizing process of learning

The first thing we agreed doing was getting to know Julia for the research. We used a lot of time solving problems in Julia and trying to figure out how to get a view of what was to be done and how to organize it. We mostly communicated through facebook and created a repository on github to share project files. We gathered our solutions to Project Euler and made comparisons. Gathering information about benchmarking we mostly ran into walls of text on how to do it properly, and did not know where to start. Every time we thought we were understanding benchmarking, our supervisor taught us something new. We created a board on Trello. We made a to-do list of cards so we knew what had to be done and started handing out tasks. We communicated more with our supervisor via email as we started writing the report. The group was really good at getting the tasks, we agreed on doing, done. We should have been using more tools for organizing and we did not set any deadlines for anything. Scrum is a tool, which would have helped us a lot, but it is not that easy to learn. We should have met and worked together, several times a week, discussing our problems and goals.

14.1.2 Process of learning

We learn by both direct and indirect learning process. We learned Julia indirectly by solving Project Euler. We pretty much read about every topic, which is direct learning. We cannot think of another way to learn programming than reading, trying it out and reading to get a result. Project Euler was a good support for learning, since learning something new, is much easier, when there's a purpose or a goal. Programming is not as fun if there's no end goal in mind.

14.1.3 Cooperation, difficulties in group and with supervisor

We actually had a good idea about who was good at what and splitting the tasks between us with this in mind. The communication could have been better. It seemed that we worked on different times of the day, which made

communication difficult. Our supervisor did his best and we think he did a good job.

14.1.4 Conclusion

Trello is a very good tool for organizing and we should have been using it from the beginning. Looking beyond the fact that some data was lost using Git, the tool is perfect for team programming. We have payed the price not making any deadlines. Deadlines seems useless when the focus is just on the project, but when the workload gets more intense, it can be very good to have smaller deadlines or milestones during a project. Meetings should have occured more often, and syncing the current progress. Overall collaboration was very well, taken into account what could have been improved.

14.2 Copy of poster in A4

References

- [1] Thomas H. Cormen, Charles E. Leiserson, Ronald Rivest, Clifford Stein *Introduction to Algorithms*, Cambridge, Massachusetts, third edition, 2009.
- [2] http://www.tiobe.com/tiobe_index
- [3] <http://julialang.org/blog/2012/02/why-we-created-julia>
- [4] <http://docs.julialang.org/en/latest/manual/parallel-computing/>
- [5] <http://unicode.org/standard/WhatIsUnicode.html>
- [6] <http://www.cs.cornell.edu/courses/cs1112/2013fa/Exams/exam2/vectorizedCode.pdf>
- [7] <https://projecteuler.net/>
- [8] <http://dictionary.cambridge.org/dictionary/english/benchmark>
- [9] <http://dictionary.cambridge.org/dictionary/english/benchmark>